

# Performance Provisioning in MapReduce



Reza Dibaj

# Contents:

- History of Databases
- NoSQL
- Hadoop
- Structured vs Unstructured
- Relational Tables vs Key-Value
- MapReduce
  - Introduction
  - Architecture
  - Basic Concepts
  - DBMS vs MapReduce
  - Advantages and Pitfalls

# History

**1980s**

**SQL, a dominant database environment**




Data Persistency

Integration and Reporting for many Organizations



Impedance mismatch



# History

**1990s**

## **Rise of Object Databases**

- In memory structure → Save data directly from memory to disk

Still we have Relational Dominance

# History

2000s

Amazon  
Google



High Traffic

limitation scaling up, we will hit a ceiling on RAM/CPU we can use on larger and larger boxes. so we need to resort to soln 2

1<sup>st</sup> Solution: Scale up the hardware  
Buy Bigger Boxes (BBB)

2<sup>nd</sup> Solution: Scale out the hardware  
Buy Lots of Little Boxes "commodity hardware"



**Problem:** SQL was designed to work on a large box, and not on a large amount of small boxes.

# History

## 2000s    NoSQL Movement

Some organizations started to work on something rather than traditional SQL environments

Google → BigTable (Column Family)    sth similar to excel

Amazon → Dynamo (Key-Value)

regulations to write data for RDBMS  
regulations to READ data for NoSQL


# NoSQL

Whenever we talk about Big Data we talk about unstructured data.

## Some NoSQL Characteristics

- Non-relational (Unstructured)
- Cluster-friendly referring to the commodity hardware limitation we discussed prev
- Open-source

### *HOW TO WRITE A CV*



# NoSQL

## NoSQL Data Models

- Key-value
  - Document data model
  - Column-family
  - Graph
- Aggregate-Oriented, Not ACID
- Graphs are ACID (Atomic, Consistent, Isolated, Durable)

## All NoSQL Data Models are Schema-less

Aggregate-Oriented DBs are naturally fit-in storing data on large clusters.

Graphs are good for handling relationship between things. Even relational Databases work with a high degree of complexity with relationships, while graphs are very good at this feature. There is a kind of query language specifically designed for that.



# NoSQL

## NoSQL

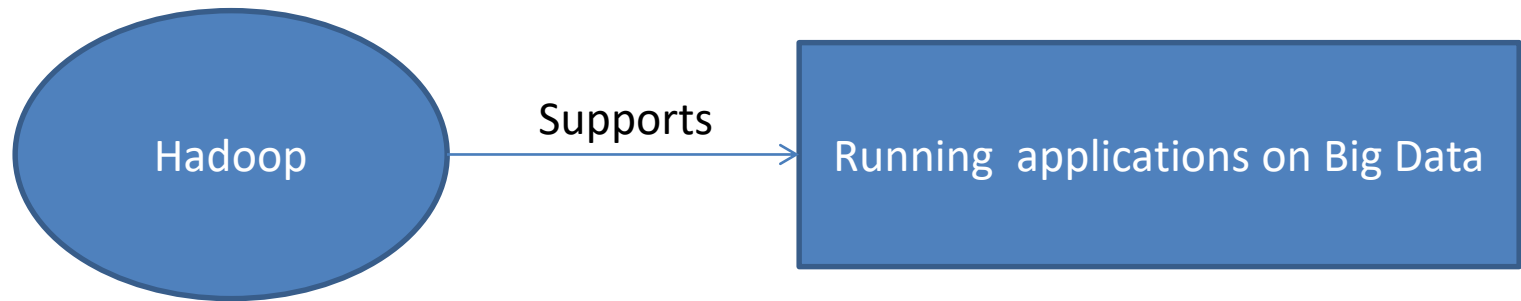
- Easier development
- Large-scale data (Big Data)

## Is NoSQL the future of Databases?


We will have different models of Databases, and we should chose the appropriate one for our specific problem.

will not replace RDBMS, they have their own uses

# Hadoop Objectives




# Open-Source Environment



# Big Data Challenge Points

The 7's of Big Data: <https://impact.com/marketing-intelligence/7-vs-big-data/>




variability, veracity, visualization, value

visualization: tool to bring information to knowledge

variability:


# Traditional Approach

Enterprise Approach




# Traditional Approach

Enterprise Approach



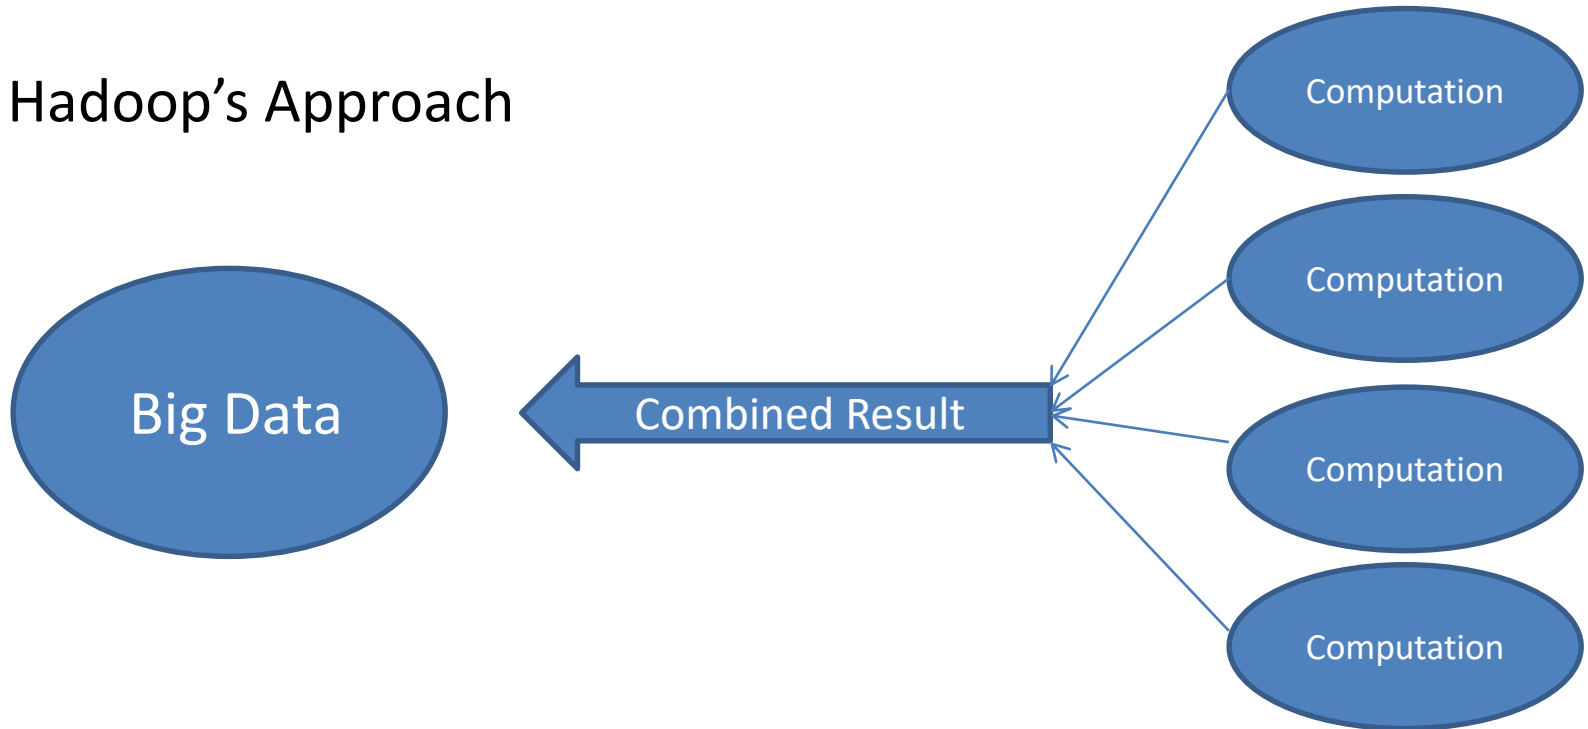
# Breaking the Data

Hadoop's Approach




# Move Computation to the Data

Hadoop's Approach



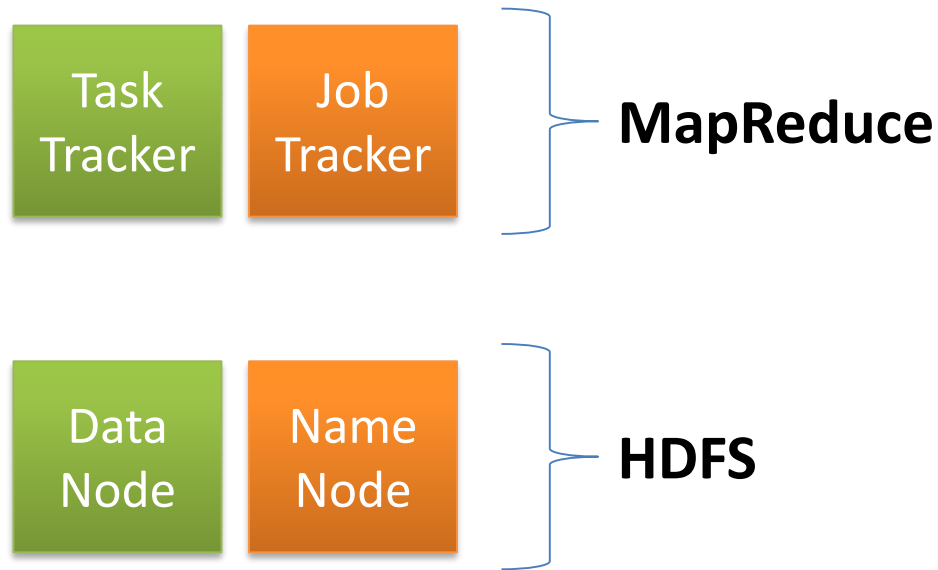


# Hadoop's Architecture



# MapReduce and HDFS

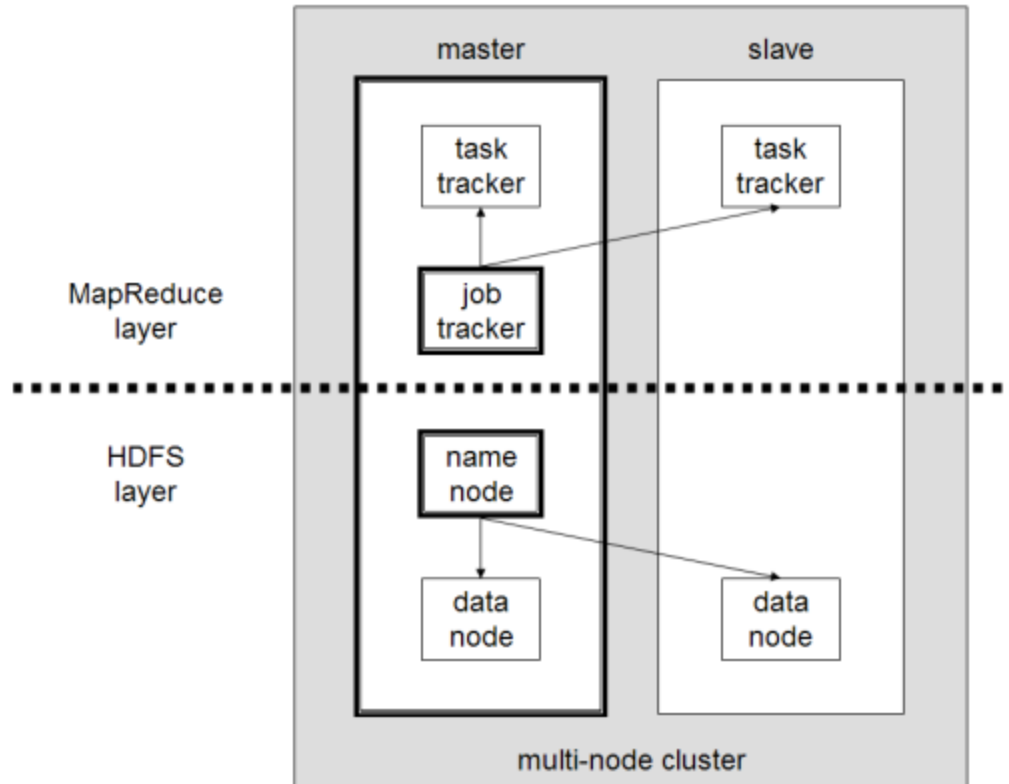
job split into smaller tasks



**Any application contacts the Master Node.  
Then the task will go into the queue.**

# Hadoop's Architecture


master splits job into tasks for slaves; it can also take on role of slave and do work too in its free time after delegating the work and before collecting results from slaves



# Distributed Model


Low Cost  
Computers

Works on  
Linux Based  
Machines




# Task Trackers and Data Nodes

**Slaves**



# Master




**Master**

splitting job into tasks and the I/O related to this process is the most time consuming task in this master -> slave process

this I/O is 200x slower than when compared to the I/O in a central model

however, after the I/O bottleneck the performance goes through the roof :)


**Slaves**



# MapReduce and HDFS

An orange square icon with the text "Job Tracker" in white.


Job  
Tracker

A blue bracket pointing from the Job Tracker icon to the text.

Break higher bigger task to smaller pieces and to send each smaller piece of computation to the task trackers. It also can combine the results and send it back to the application.

An orange square icon with the text "Name Node" in white.

Name  
Node

A blue bracket pointing from the Name Node icon to the text.

It is responsible to keep the index that which data is resides on which data node. When application contacts a Name Node, it tells the application that in order to get your particular data; you should go to which computer.

# An Introduction to MapReduce

## MapReduce



- Simple
- Scalable
- Fault-tolerant



- limitations on its performance and efficiency

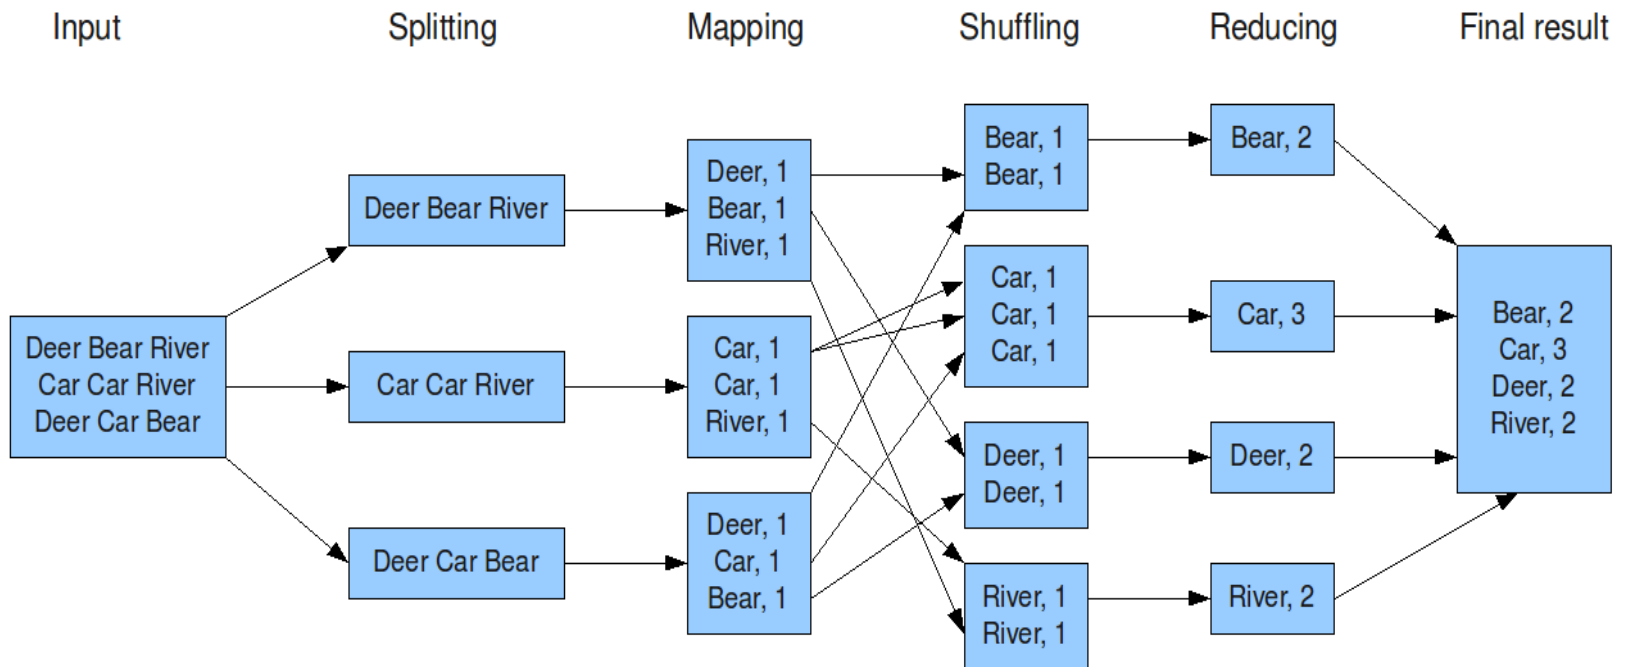
this is why we will not replace RDBMS:

1. performance
2. security








# MapReduce Architecture

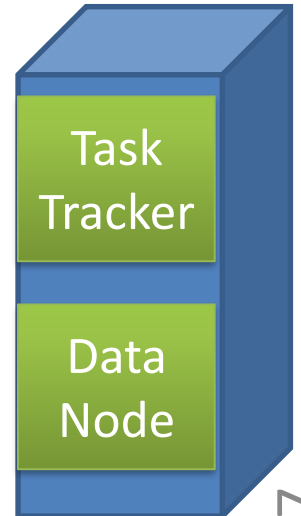




The overall MapReduce word count process




# Fault Tolerance for Data



# Fault Tolerance for Processing



# Scalable



# Easy Programming

Programmers



Do not have to worry about

Where the program is located

How to manage the failures

How to break computations into smaller pieces

How to program for scaling

# Easy Programming


Programmers



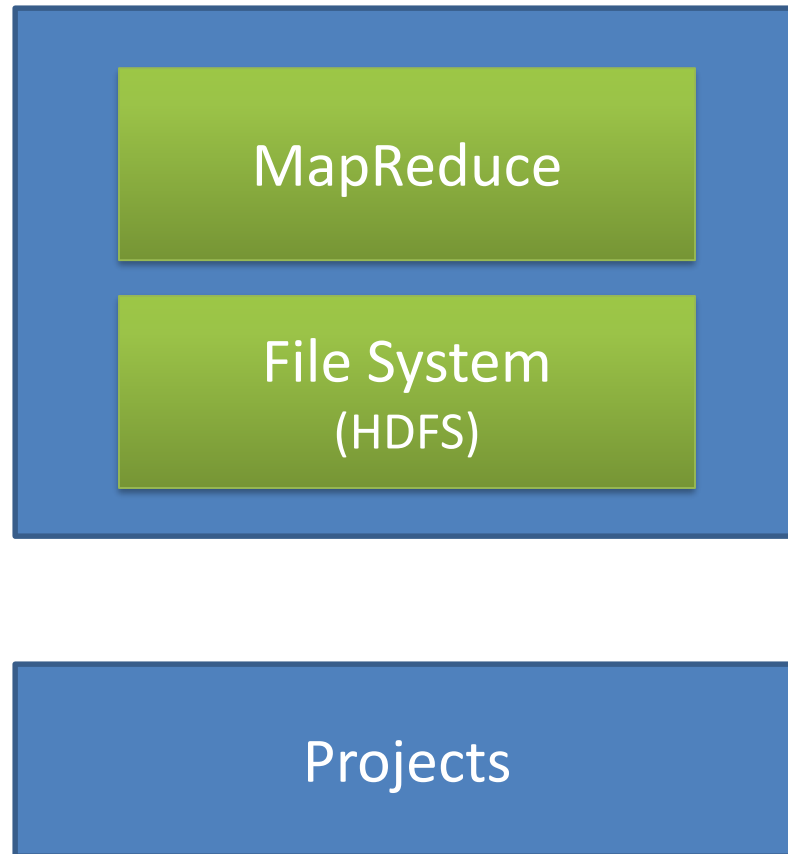
Could focus on

Writing Scale-free Programs

# Scalability Cost

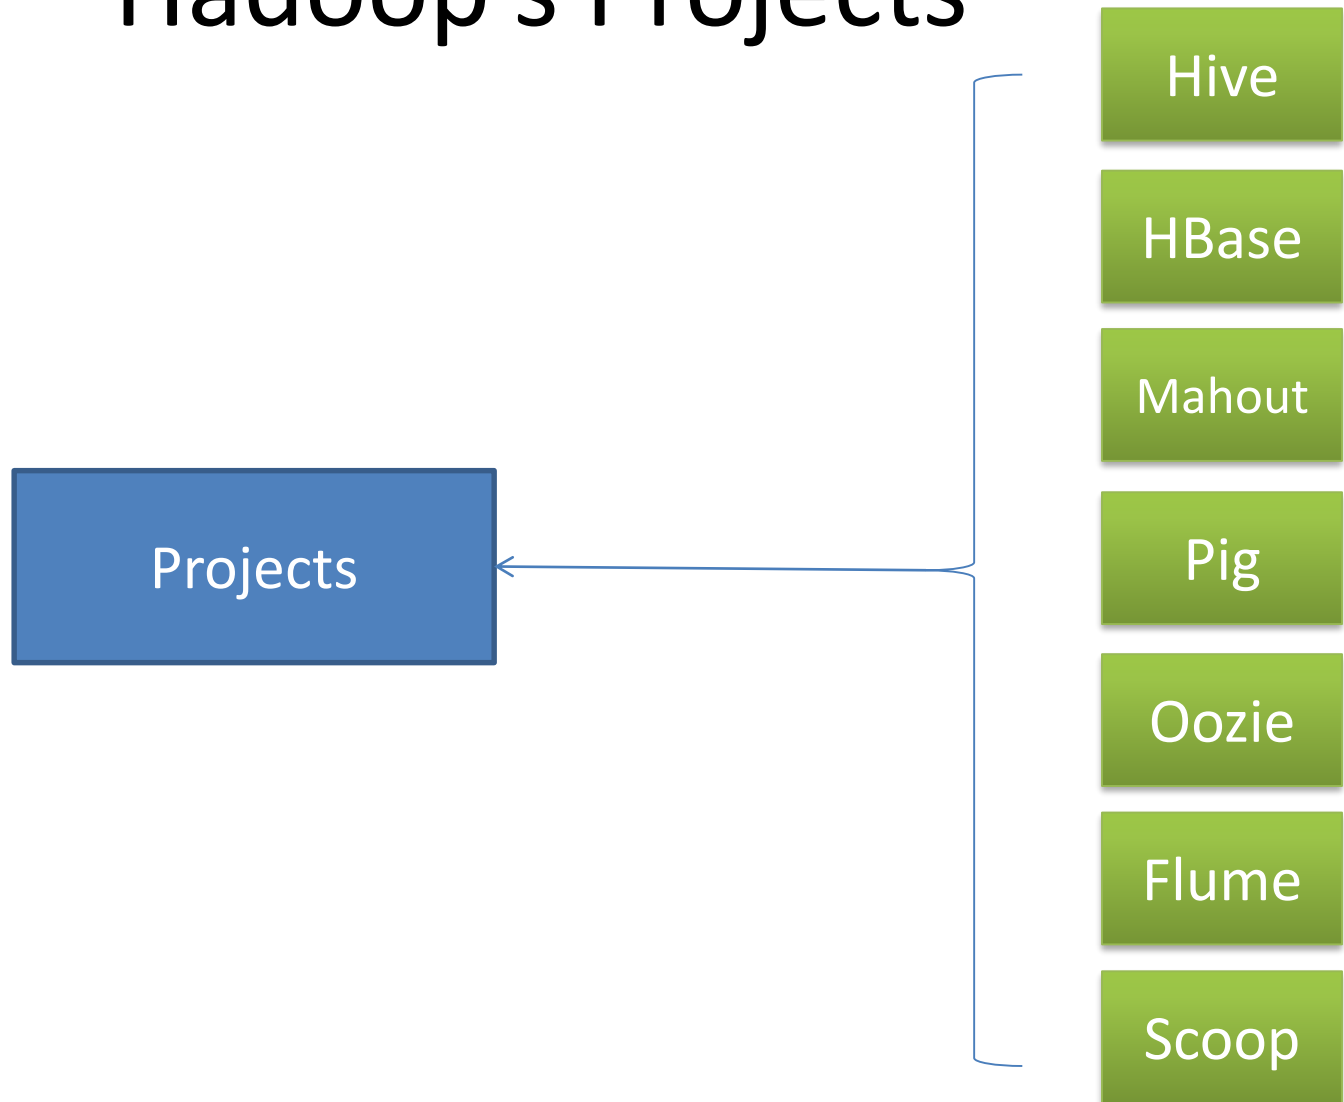


# Hadoop's Projects





# Hadoop's Projects



# Little about Hadoop's History

2004-2005

Hadoop was created in 2005 by Doug Cutting and Michael Cafarella who both worked for Yahoo!

2006

In 2006 the project was donated to Apache.

# Hadoop's Users

## Administrators

- Installation
- monitoring/managing the system
- tuning the system

## Users

- Design Applications
- Import/Export Data
- Work with tools

# Hadoop's Usage Areas

- Social Media
- Retail
- Financial Services
- Searching Tools
- Governments
- Intelligence

# Hadoop's User Companies

- Yahoo
- Facebook
- Amazon
- eBay
- American Airline
- The New York Times
- Federal Reserve Board
- Chevron
- IBM

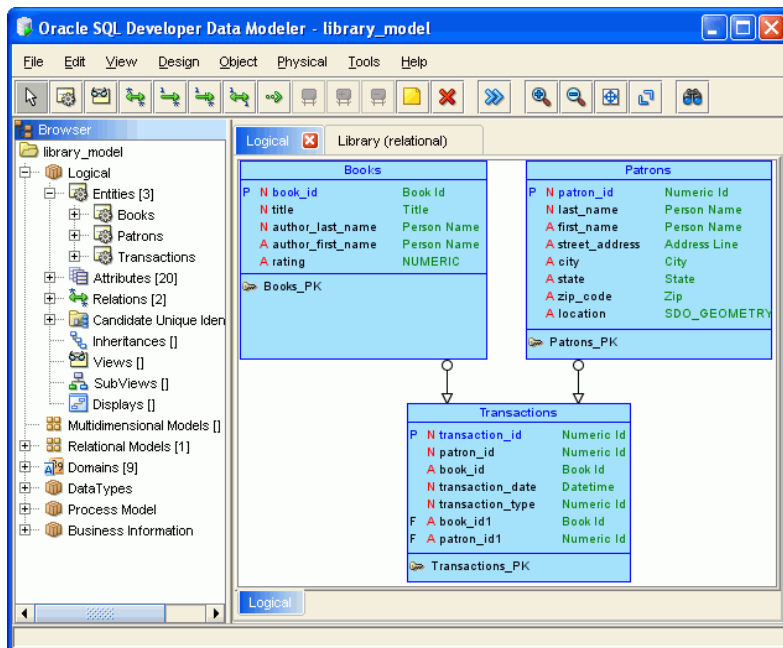
# Outlook of Hadoop

- By 2015 more than 50% of Enterprise Data was processed by Hadoop

# Structured vs Unstructured


## SQL

is by design targeted at structured data.



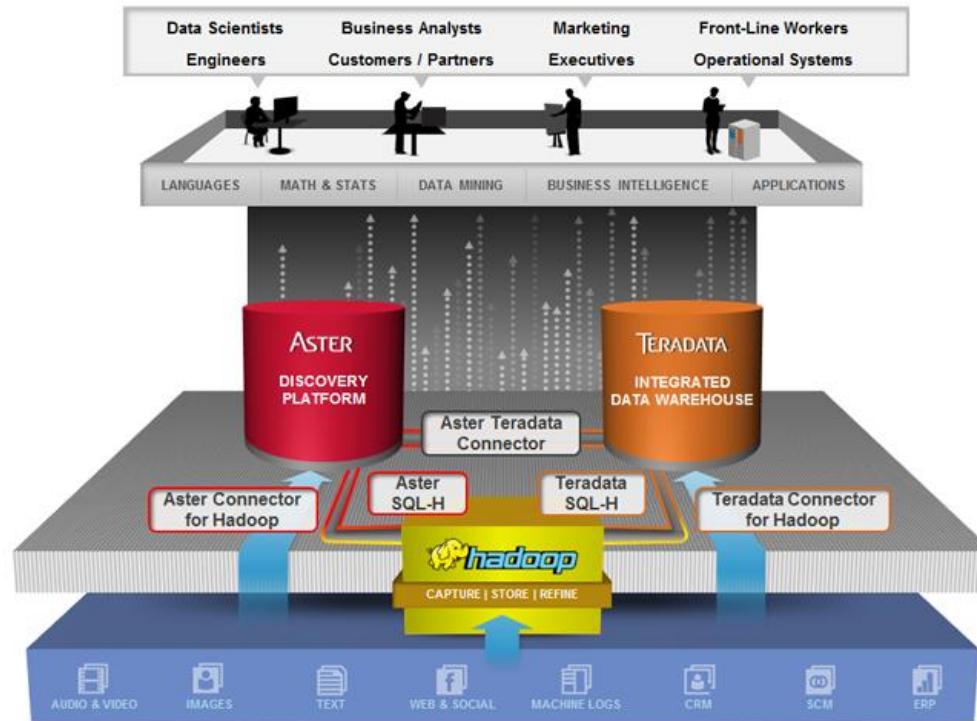
## Hadoop

many of its initial applications deal with unstructured data such as text.



# Structured vs Unstructured

SQL and Hadoop can be complementary, as SQL is a query language which can be implemented on top of Hadoop as the execution engine.

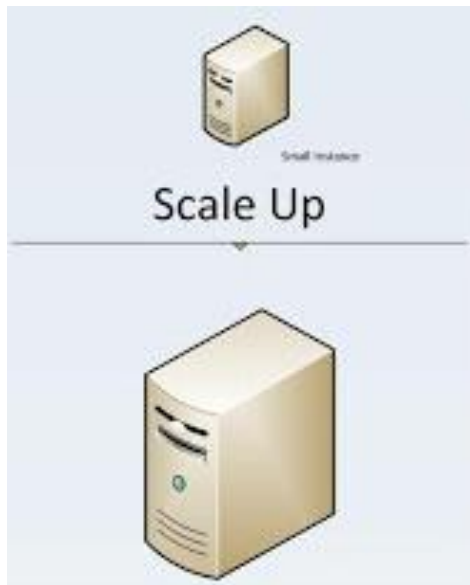


can have hive, pig, etc. on top of hadoop




# Structured vs Unstructured

**SQL** design is more friendly to **scaling up**.

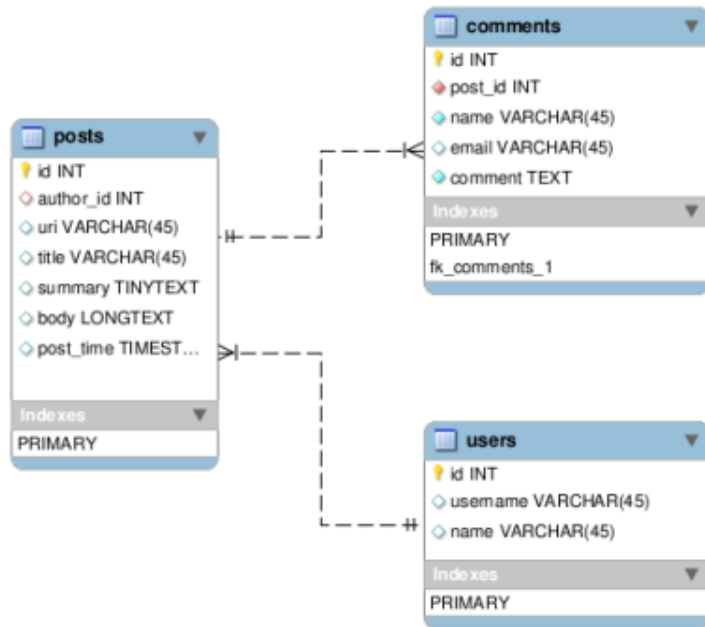


**Hadoop** is designed to be a **scale-out** architecture operating on a cluster of commodity PC machines.




# Relational Tables vs Key-Value

In **SQL** data resides in tables having relational structure defined by a schema.



**Hadoop** uses key/value pairs as its basic data unit, which is flexible enough to work with the less-structured data types.



# Declarative Query vs Functional Programming

Under **SQL** you have query statements.

```
SELECT Book.title AS Title, COUNT(*) AS  
Authors  
FROM Book  
JOIN Book_author  
ON Book.isbn = Book_author.isbn  
GROUP BY Book.title;
```

Under Hadoop you have scripts and codes.  
MapReduce allows you to process data in a more general fashion than SQL queries.

```
function map(String name, String document):  
  // name: document name  
  // document: document contents  
  for each word w in document:  
    emit (w, 1)  
  
function reduce(String word, Iterator partialCounts):  
  // word: a word  
  // partialCounts: a list of aggregated partial counts  
  sum = 0  
  for each pc in partialCounts:  
    sum += ParseInt(pc)  
  emit (word, sum)
```

# DBMS vs MapReduce

## DBMS

- Generates a query plan tree for execution

## MapReduce

- A plan for executions in MapReduce is determined entirely at runtime

# Pros and Cons

## Debate

- DBMSs have adopted “one size fits all” strategy and are not suited for solving extremely large scale data processing tasks.
- MapReduce is referred to as a new way of processing big data in data-center computing.

## Slow

- Hadoop is 2~50 times slower than parallel DBMS except in the case of data loading.

## Poor Efficiency

- Hadoop system is scalable, but achieves very low efficiency per node, less than 5MB/s processing rates, repeating a mistake that previous studies on high performance systems often made by “focusing on scalability but missing efficiency”
- Loosing efficiency for the fault tolerance: MapReduce increases the fault tolerance of long-time analysis by frequent checkpoints of completed tasks and data replication.

# MapReduce Advantages and Pitfalls



- Simple and easy to use
- Flexible
- Independent of the storage
- Fault-tolerance
- High scalability



- No high-level language
- No schema and no index
- A single fixed dataflow
- Low efficiency
- Very young

indexes are for making processing faster

hadoop has sth like indexes diff from rdbms

# References

- [1]Lee, K. H., Lee, Y. J., Choi, H., Chung, Y. D., & Moon, B. (2012). Parallel data processing with MapReduce: a survey. *AcM SIGMoD Record*, 40(4), 11-20. doi: 10.1145/2094114.2094118
- [2]Pavlo, A., Paulson, E., Rasin, A., Abadi, D. J., DeWitt, D. J., Madden, S., & Stonebraker, M. (2009, June). A comparison of approaches to large-scale data analysis. In *Proceedings of the 2009 ACM SIGMOD International Conference on Management of data* (pp. 165-178). ACM.
- (3) The performance of MapReduce: an in-depth study. doi: 10.1145/1559845.1559865
- [3]Jiang, D., Ooi, B. C., Shi, L., & Wu, S. (2010). The performance of MapReduce: an in-depth study. *Proceedings of the VLDB Endowment*, 3(1-2), 472-483. doi: 10.14778/1920841.1920903
- [4]Eltabakh, M. Y., Özcan, F., Sismanis, Y., Haas, P. J., Pirahesh, H., & Vondrak, J. (2013, March). Eagle-eyed elephant: split-oriented indexing in hadoop. In *Proceedings of the 16th International Conference on Extending Database Technology* (pp. 89-100). ACM. doi: 10.1145/2452376.2452388
- [5]Chen, G., Vo, H. T., Wu, S., Ooi, B. C., & Özsu, M. T. (2011). A Framework for Supporting DBMS-like Indexes in the Cloud. *Proceedings of the VLDB Endowment*, 4(11), 702-713. doi: 10.14778/2556549.2556550

*Thank you*

