

Naïve Bayes

Following Applications:

- Spam Classification – Given an email, predict whether it is spam or not .
(Spam or Ham)
- Medical Diagnosis – Given a list of symptoms, predict whether a patient has disease X or not
- Weather – Based on temperature, humidity, etc... predict if it will rain tomorrow

Characteristics

- The relationship between attribute set and the class variable is non-deterministic.
- Even if the attributes are same, the class label may differ in training set even and hence can not be predicted with certainty.
- Reason: noisy data, certain other attributes are not included in the data.

Prediction with limited attributes

- Task of predicting whether a person is at risk for heart disease based on the person's diet and workout frequency.
- So need an approach to model probabilistic relationship between attribute set and the class variable.

Problem statement:


- Given features X_1, X_2, \dots, X_n – Predict a label Y

Bayesian Classifiers

- compute the posterior probability $P(C \mid A_1, A_2, \dots, A_n)$ for all values of C using the Bayes theorem
- Choose value of C that maximizes $P(C \mid A_1, A_2, \dots, A_n)$

Another Application

- Digit Recognition




- $X_1, \dots, X_n \in \{0, 1\}$ (Black vs. White pixels)
- $Y \in \{5, 6\}$ (predict whether a digit is a 5 or a 6)
- Cycle thru all the image cells (pixels) predicting a 5 or 6
- Compute $P(5, X_1, \dots, X_n)$ and $P(6, X_1, \dots, X_n)$ and decide if it is 5 or 6.

Naiive Bayes Training

Naïve Bayes Training

Now that we've decided to use a Naïve Bayes classifier, we need to train it with some data:




MNIST Training Data






Naïve Bayes Training

- For binary digits, training amounts to averaging all of the training fives together and all of the training sixes together.



Naiive Bayes Classification



Bayes Theorem

- Abstractedly stated:

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}.$$

- Stated as conditional probability:
- $(B|A)$ – can be interchanged as hypothesis and event.
- It allows us to invert the probability $(B|A)$ to $(A|B)$.

Bayes Cont'd

- $(A|B)$ – B is the event , A is the hypothesis.
- B – you have tested positive for a disease
- A – Do you have the disease.
- Inversion:
- $(B|A)$ – A you have the disease
- B – Did you test positive

Bayes – re-casted as inference, Prior and Posterior

$$P(\mathcal{H}|\mathcal{D}) = \frac{P(\mathcal{D}|\mathcal{H})P(\mathcal{H})}{P(\mathcal{D})}.$$

- The prior $P(\mathcal{H})$ is the probability that \mathcal{H} is true before the data is considered.
- The posterior $P(\mathcal{H}|\mathcal{D})$ is the probability that \mathcal{H} is true after the data is considered.
- The likelihood $P(\mathcal{D}|\mathcal{H})$ is the evidence about \mathcal{H} provided by the data \mathcal{D} .
- $P(\mathcal{D})$ is the total probability of the data taking into account all possible hypotheses.

Bayes and the Base rate Fallacy

- A screening test for a disease is both sensitive and specific.
- Positive when testing a person with the disease
- Negative when testing someone without the disease.
- True positive rate is 99%.
- False positive rate is 2%.
- Suppose the prevalence of the disease in the general population is 0.5%.
- If a random person tests positive, what is the probability that they have the disease?
- Most people would guess that the answer is 99%.....

Using the notation established above for hypotheses and data:

Let:

- H^+ be the hypothesis that the person has the disease
- H^- be the hypothesis they do not.
- T^+ and T^- represent the data of a positive and negative screening test respectively.
- We are asked to compute $P(H^+ | T^+)$.
- Given $P(T^+ | H^+) = 0.99$, $P(T^+ | H^-) = 0.02$, $P(H^+) = 0.005$.


Expressing these facts in a tree

- Completing the missing part
- $P(T- | H+) = 0.01$, $P(T- | H-) = 0.98$

- $P(T+ | H+) = 0.99$

$$P(T+ | H-) = 0.02$$

- $P(H+)$ $P(H-)$



Calculation by Bayes Theorem

Bayes' theorem yields

$$P(\mathcal{H}_+|\mathcal{T}_+) = \frac{P(\mathcal{T}_+|\mathcal{H}_+)P(\mathcal{H}_+)}{P(\mathcal{T}_+)} = \frac{0.99 \cdot 0.005}{0.99 \cdot 0.005 + 0.02 \cdot 0.995} = 0.19920 \approx 20\%$$

Computing with tables

- Construct the table as follows. Pick a number, say 10000 people, and place it as the grand total in the lower right. Using $P(D+) = .005$ we compute that 50 out of the 10000 people are sick (D+).
- Likewise 9950 people are healthy (D-). At this point the table looks like:

	$D+$	$D-$	total
$T+$			
$T-$			
total	50	9950	10000

Computing with tables....

- Using $P(T+|D+) = .9$ we can compute that the number of sick people who tested positive as 90% of 50 or 45. The other entries are similar. At this point the table looks like the table below on the left.
- Finally we sum the T+ and T- rows to get the completed table on the right.

	<i>D+</i>	<i>D-</i>	total
<i>T+</i>	45	498	
<i>T-</i>	5	9452	
total	50	9950	10000

	<i>D+</i>	<i>D-</i>	total
<i>T+</i>	45	498	543
<i>T-</i>	5	9452	9457
total	50	9950	10000

Computing with tables....

Using the complete table we can compute

$$P(D+|T+) = \frac{|D+ \cap T+|}{|T+|} = \frac{45}{543} = 8.3\%.$$

Computing with tables

Symbols: For completeness, we show how the solution looks when written out directly in symbols.

$$\begin{aligned}P(D+|T+) &= \frac{P(T+|D+) \cdot P(D+)}{P(T+)} \\&= \frac{P(T+|D+) \cdot P(D+)}{P(T+|D+) \cdot P(D+) + P(T+|D-) \cdot P(D-)} \\&= \frac{.9 \times .005}{.9 \times .005 + .05 \times .995} \\&= 8.3\%\end{aligned}$$

Naïve Bayes - Acknowledgement

- B S Panda – IIT (Delhi)
- Jeremy Orloff and Jonathan Bloom - MIT